Vasilen Kanchev

Data Engineer

Porto, Portugal (remote) | EU citizen (Bulgarian), no visa needed in the EU | +351 931 673 151 | vasilen.business@gmail.com
LinkedIn | GitHub | Portfolio

WORK EXPERIENCE

Data Engineer and Chief Revenue Officer | Gosian Media LLC
South Carolina | June 2025 - Present

- Lead 80 employees. Personally recruited and trained 30+ of them and wrote the SOPs and training manuals that standardize onboarding.
- Built a revenue-versus-cost tracking system by department, then cross-trained flexible staff to cover two roles: cut cost of goods sold by about $10K a month and raised performance in the strongest departments by KPI.
- Joined at about $80K in monthly revenue; now about $600K a month (7.5x), with profit up $60K+ a month after all expenses.
- Built the PostgreSQL database on AWS RDS (30+ spreadsheets migrated) and the Docker, Python and dbt pipelines with validation tests behind it, replacing 10+ hours a week of manual data entry.

Data Analyst and Senior Supervisor | Influencer Marketing Hub
United Kingdom | November 2023 - June 2025

- Led the afternoon shift of a three-shift operation, directly overseeing 150 employees and generating about $3M a month in revenue for that shift.
- Automated KPI reporting with Python and Tableau for 15+ stakeholders; built ETL from social media APIs; supported campaign A/B tests.

Data Administrator / Operations Analyst | TapIn Platforms, Inc.
Miami, Florida | Full-time | October 2021 - November 2023

- Managed relational databases, performing daily validation and record updates to support accurate operational data.
- Used advanced Excel and SQL to identify business trends, track user growth and report on operational metrics.
- Identified inefficiencies in raw data collection and proposed software solutions that centralized the company's internal knowledge base.
- Supported product and marketing teams with ad hoc data requests and SQL queries to extract specific user cohorts.

Operations & Data Assistant | CoEfficient & NPC
Johannesburg, South Africa | Part-time | July 2019 - September 2021

- Digitized legacy records and ran monthly reconciliations across CRM, inventory and finance systems, over 40 tables.

PROJECTS

E-Commerce Finance Mart | DuckDB, SQL, BigQuery, Databricks, Excel
1.07M real invoice lines (public UK retail data) in a star schema reconciled to the penny, also loaded and reconciled in BigQuery and Databricks Free Edition: 14 blocking checks, 32 tests, Excel model, explorer and Tableau dashboard. Found 44% of credit-note value was same-hour reversals. vasilendata.github.io/ecommerce-finance-mart/explorer/ | Tableau Public dashboard

A/B Test Analysis | Python, SciPy, pytest, GitHub Actions
Real 90K-player experiment with a tested toolkit (sample-ratio check, z-test, Holm, power, A/A), plus sequential testing and CUPED validated by simulation (25 tests). Gate change cut day-7 retention 0.82 points (p = 0.0016). github.com/VasilenData/ab-test-analysis

BigQuery Cost Lab | Python, SQL, BigQuery, pytest
Measured query cost on a public table: SELECT * with LIMIT would scan 2.38 TB, a day filter 85 MB; clustering cut a scan 11.6x. Tested linter predicts bytes from column types (22 tests). github.com/VasilenData/bigquery-cost-lab

HN Pulse: Live Data Pipeline | Python, dbt, Airflow, DuckDB, GitHub Actions
Hacker News API through dbt models (27 data tests, enforced contracts) into DuckDB, with a freshness gate; failing tests block the live dashboard update. vasilendata.github.io/hn-pulse/

NYC Taxi Data Platform | Dagster, dbt, DuckDB, Docker
31.6M trips with an enforced data contract; every rejected row (2.8%) kept in a quarantine table with its reason; blocking checks stop a corrupted month. 10 months build in 62 s. github.com/VasilenData/nyc-taxi-data-platform

Data Contracts Platform | Python, DuckDB, YAML, pytest
YAML contracts, DuckDB validator and breaking-change detector. Fitted on Jan 2023 NYC taxi data (3.07M rows), it failed all 11 later months on real drift (195 tests). github.com/VasilenData/data-contracts-platform

Four-Engine Benchmark on 79M Rows | DuckDB, Polars, pandas, Spark
Same four queries on 79 million NYC taxi trips with cross-checked answers: DuckDB 1.4 s against 32.8 s for pandas. github.com/VasilenData/taxi-engines-benchmark

TECHNICAL SKILLS

Python, SQL (PostgreSQL, window functions), Excel, dimensional modeling, ELT, data contracts, data quality testing, Bash, AWS RDS, Docker, dbt, Apache Airflow, DuckDB, Kafka, Spark, Polars, Terraform, Dagster, GitHub Actions, Tableau, Pandas

EDUCATION & CERTIFICATIONS

University of Suffolk | BSc (Hons) Computer Science | Expected 2027
UNISA | Mathematics and Statistics Higher Certificate (Distinction) | 2024
Certifications: Google Data Analytics Professional Certificate; IBM Data Engineering Certificate